From *what changed* to *why.*

I designed Stellar, Pixis's conversational AI analyst. Marketers ask about their campaigns in plain language and get answers they can verify, act on and share.

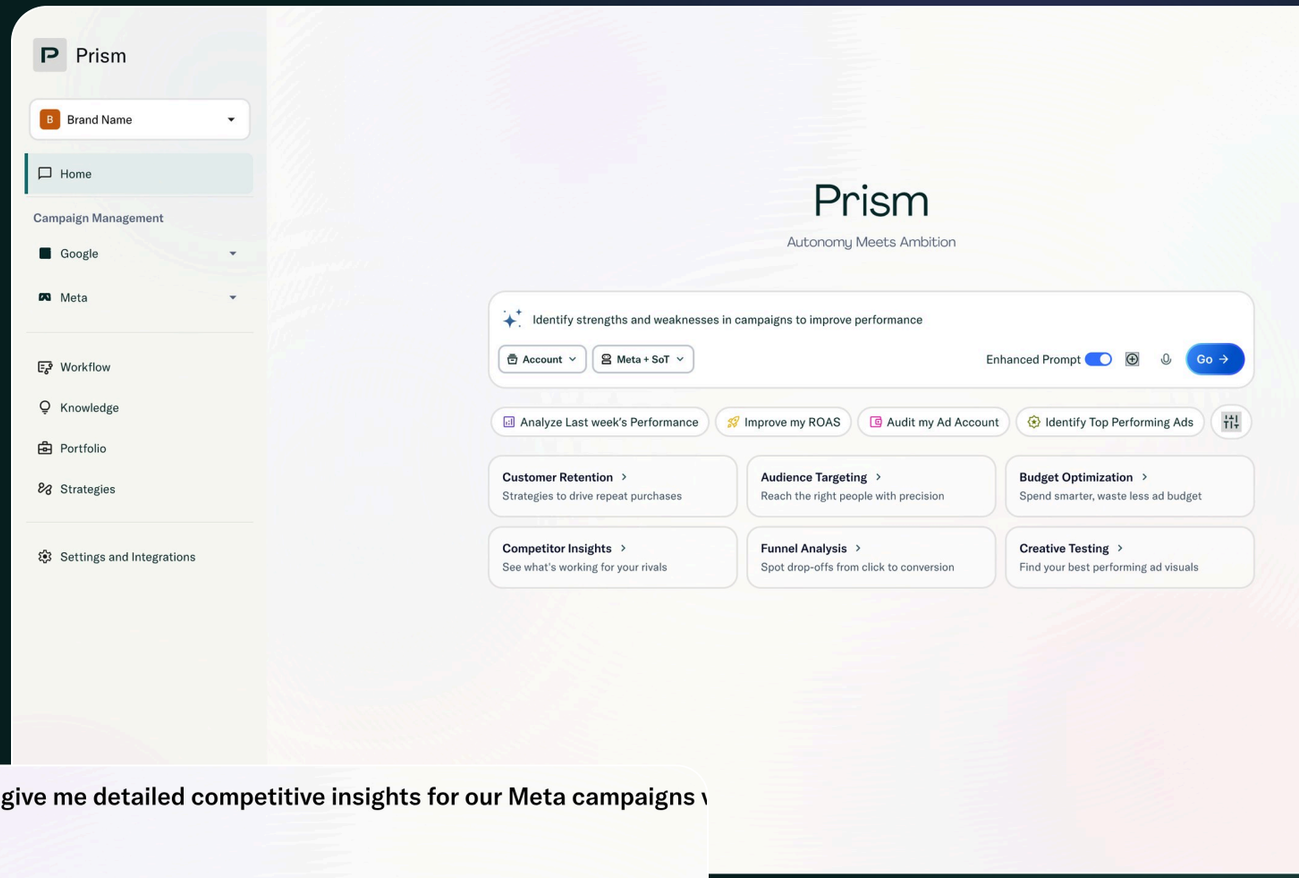
ROLE
Senior Product Designer

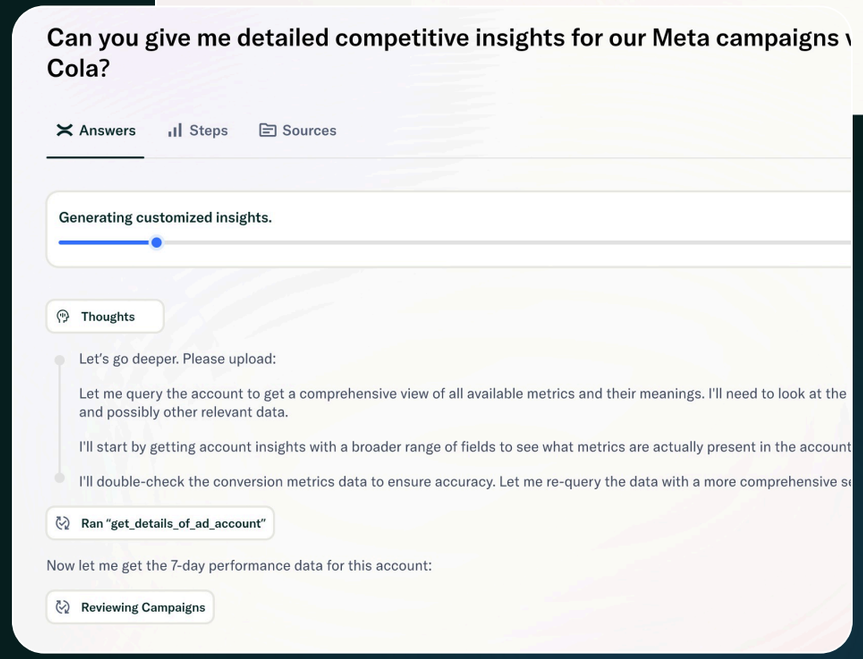
TEAM
4 product designers

SCOPE
Research, UX strategy, UI, design system, handoff

PLATFORM
Web app, desktop-first

Pixis Advance → Prism → Stellar





01 THE PROBLEM

Marketers had dashboards. What they didn't have was the *why.*

PROBLEM STATEMENT

Agency marketers can see that an account dipped, but explaining *why*, and what to do next, means digging through platforms, exports and colleagues. Their tools flag problems and stop there.

How might we let them ask the question directly, and trust the answer enough to put it in front of a client?

WHO I DESIGNED FOR

Specialists · Strategists · Account managers · Channel leads

Each one runs 2–20 client accounts across Google, Meta, Bing and programmatic platforms.

WHERE THE HOURS WENT

Finding a pattern in exported data — **3–4 h**

SENIOR SPECIALIST · PER ANALYSIS

Keyword analysis and report prep — **2 h**

PAID SEARCH MANAGER · EVERY DAY

Root cause, after a tool flags an issue — **30 min**

STRATEGIST OVERSEEING 15–20 ACCOUNTS · PER ISSUE

"It takes time to dig into the insights, write down why that happened and the next steps."

PAID SEARCH MANAGER

2–20 accounts per person

7 tabs open in a single session

4+ ad platforms per agency

40+ tools in the mix

02 RESEARCH

I sat inside the workflow before designing a single screen.

8
practitioners interviewed and observed

2
digital agencies, plus Pixis account leads

40+
tools mapped across their day

60min
per recorded session



WHO WE SPOKE TO

- **3** account management
- **3** strategy and channel leads
- **2** hands-on specialists

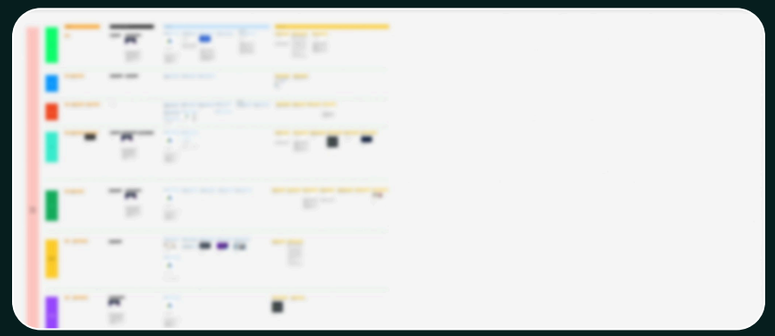
HOW I RAN IT

Interviews
Semi-structured, recorded, transcribed

Shadowing
Live working sessions, timestamped

Synthesis boards
Notes, unknowns, findings, hypotheses

Workflow maps
Team models, roles, tool walkthroughs



Workflow-mapping board, blurred on purpose: it holds participant names, agency clients and live ad-account data shared under confidentiality.

Google Ads · Meta Ads · Bing Ads · DV360 · The Trade Desk · Amazon DSP · Taboola · TikTok Ads · LinkedIn Ads · SEMrush · SpyFu · Ahrefs · Looker Studio · Datorama · Northbeam · Triple Whale · Kochava · AppsFlyer · Shopify · GA · HubSpot · Optmyzr · Feedonomics · Sheets · Excel · Slack · Asana · ClickUp · ChatGPT · + more

"Preparing and processing data is the most time-consuming part."
SENIOR SEARCH SPECIALIST

"Trusts AI recommendations, but still checks what action has been taken."
FROM MY INTERVIEW NOTES

Four insights that shaped every decision.

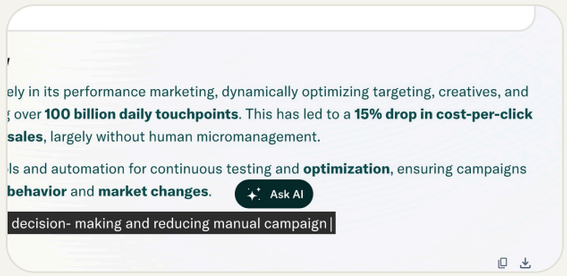
01

The bottleneck is *why*, not *what*

"[The tool] will tell you what is going wrong and where... But RCA is something that you will have to do on your own."

SO STELLAR

Leads with the explanation and suggests the next question



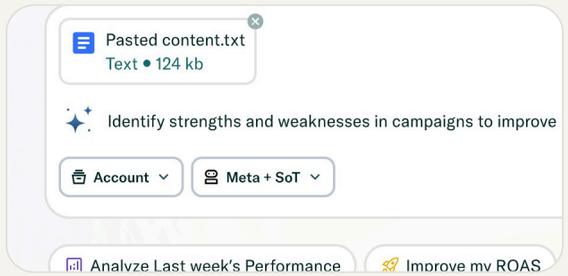
02

Every analysis is a comparison

Month over month, then week over week. Platform numbers checked against analytics or the CRM, the "source of truth."

SO STELLAR

Scopes every question to accounts and data sources up front



03

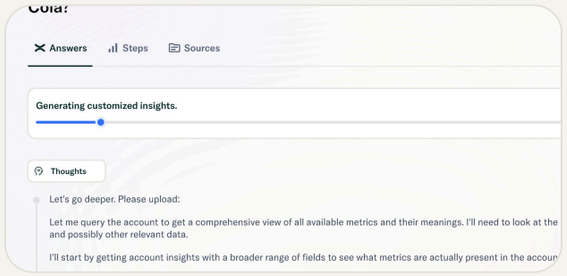
Trust comes from being checkable

"There are some times where it makes mistakes... so I can never trust it fully."

PAID SEARCH MANAGER, ON CHATGPT

SO STELLAR

Shows its steps, tools and sources, and makes errors easy to report



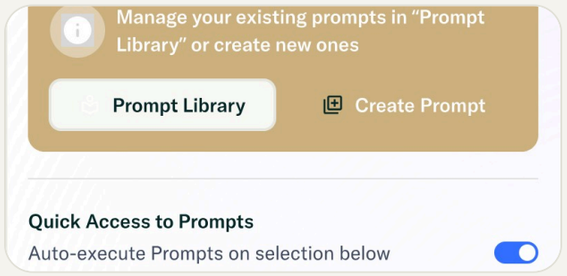
04

The work repeats, and it travels

The same morning checks run across 10–20 accounts, and the findings end up in decks, QBRs and Slack.

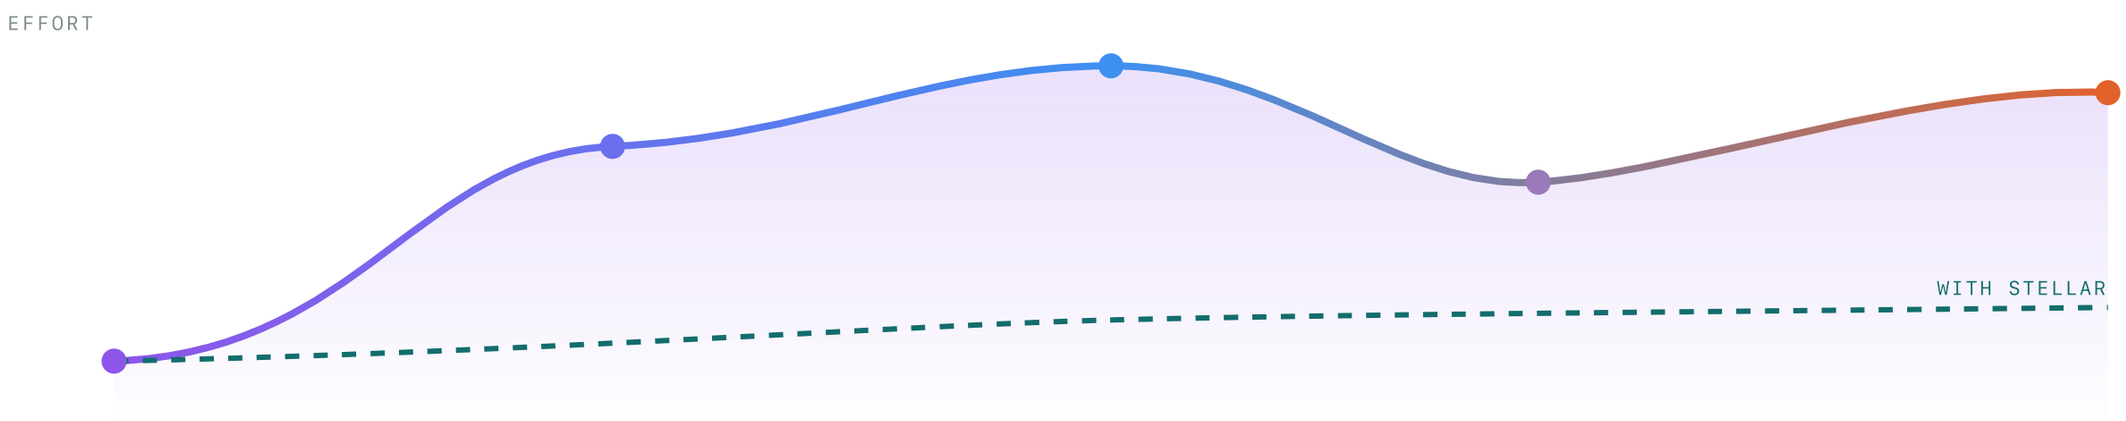
SO STELLAR

Makes prompts reusable and answers easy to share and export



04 THE JOURNEY

One strategist's evening, and where the effort piled up.



7:10 PM
Spot the dip
Sorts campaigns, scrolls row by row

7:16 PM
Dig across 7 tabs
Platform UI, Slack, issue tracker

~30 MIN
Find the root cause
MoM, then WoW, competitors, search terms

NEXT DAY
Escalate
Specialists flag issues late

WEEKLY
Report it
"I just want to make sure we have reasons, if client sees this."

THE OPPORTUNITY

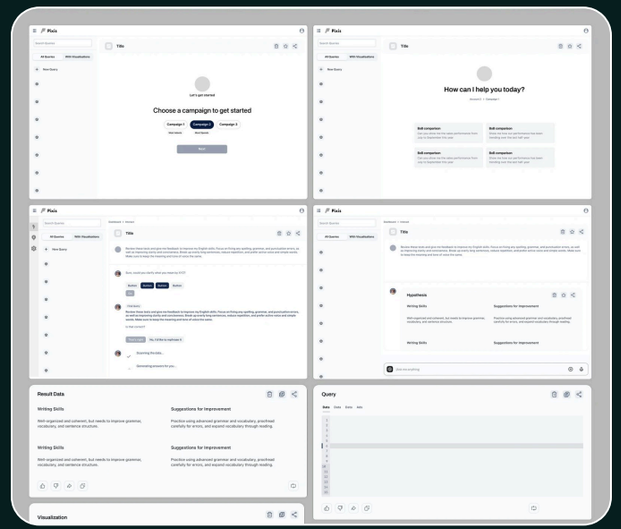
Collapse the dig, reconcile and report steps into one conversation that's grounded in their own account data.

- Answer the why, show the how
- Scope before you ask
- Chat inside the product
- Reusable questions

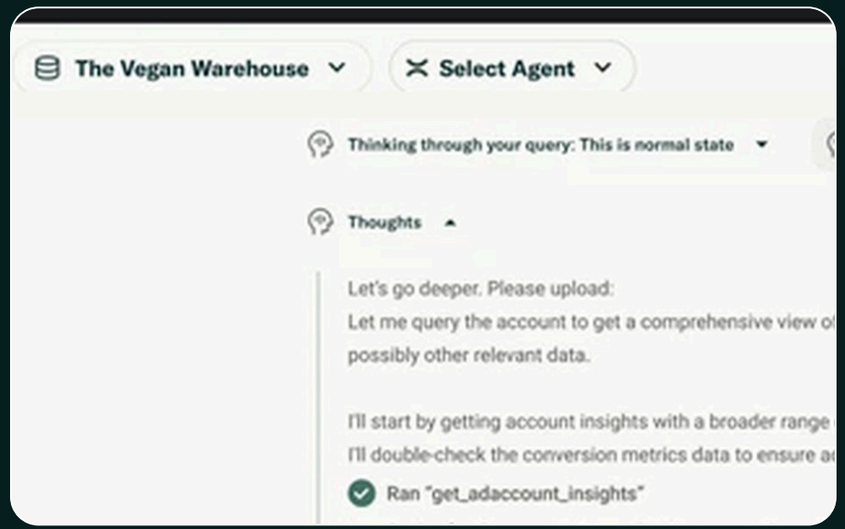
05 EVOLUTION

Three versions, three big calls.

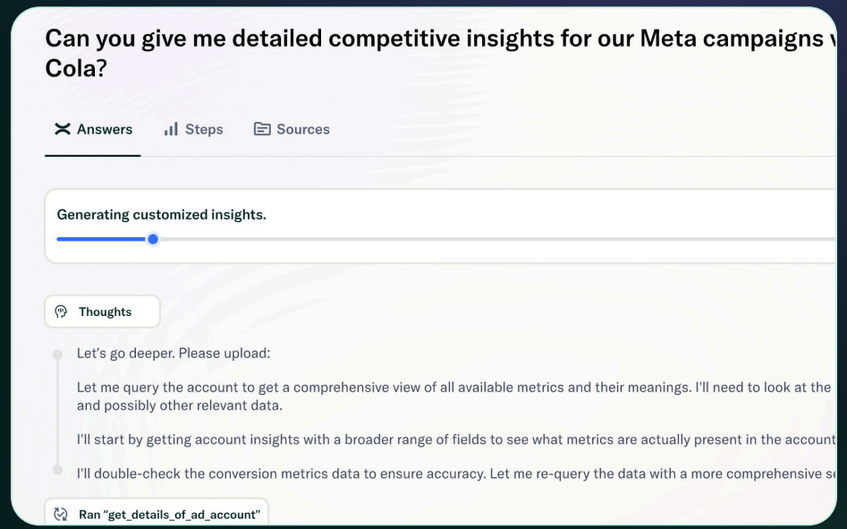
V0 · CONCEPT WIREFRAMES



V1 · PIXIS ADVANCE



V2 · PRISM → STELLAR



A Scope the data, not the agent

Advance asked users to pick an agent. Marketers think in accounts, so I swapped the picker for **Account** and **source-of-truth** chips and let the system route to the right tools.

B Separate the answer from the working

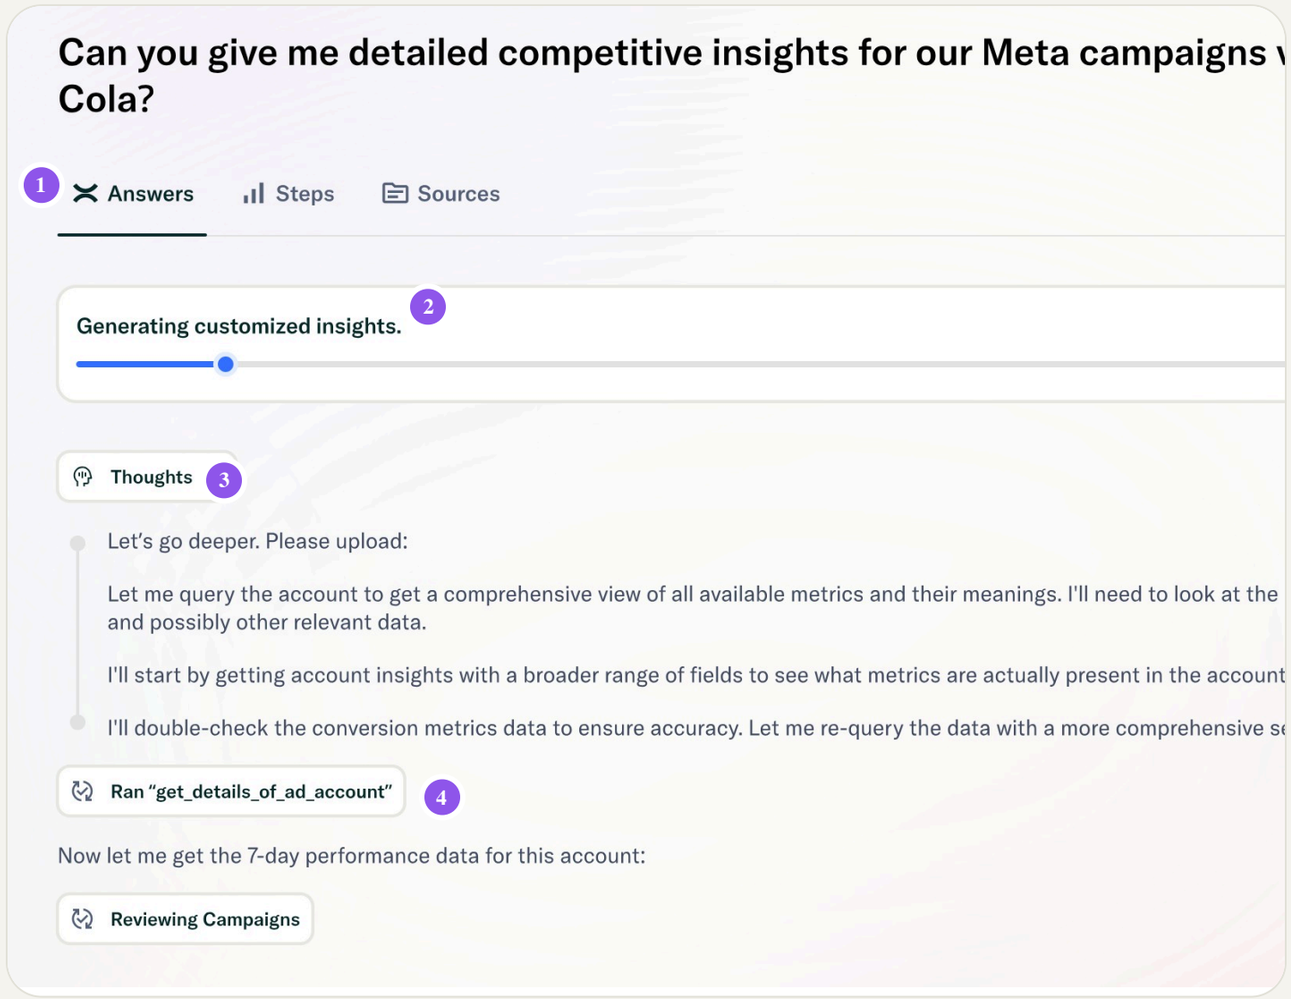
In Advance, the thoughts streamed above the answer. I split them into **Answers, Steps and Sources**, so readers see the answer first and the trail is one tap away.

C End every answer with a question

Finding the why takes several rounds, so I added **suggested next steps** and **inline Ask AI** on any highlighted sentence. Each follow-up keeps its context.

WHAT I BORROWED — **ChatGPT** → streaming, clean, short answers — **Claude** → memory and context across a thread — **Perplexity** → sources and follow-up prompts

Making a multi-step agent readable while it works.

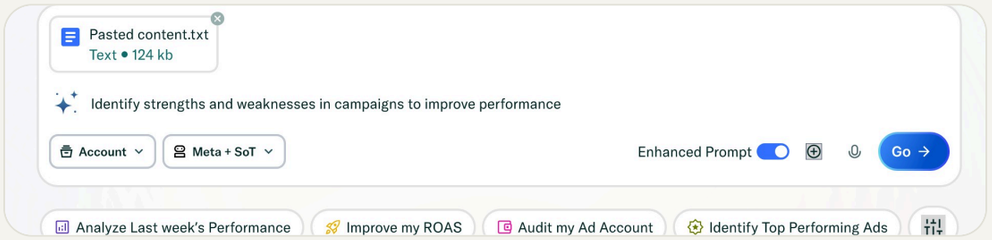


1 Answers · Steps · Sources
Answer first, trail one tap away

2 Progress in plain words
"Generating customized insights" instead of a spinner

3 Collapsible thoughts
The agent narrates its plan and says when it re-checks its data

4 Tool-call chips
You can see exactly what it touched: "get_details_of_ad_account"



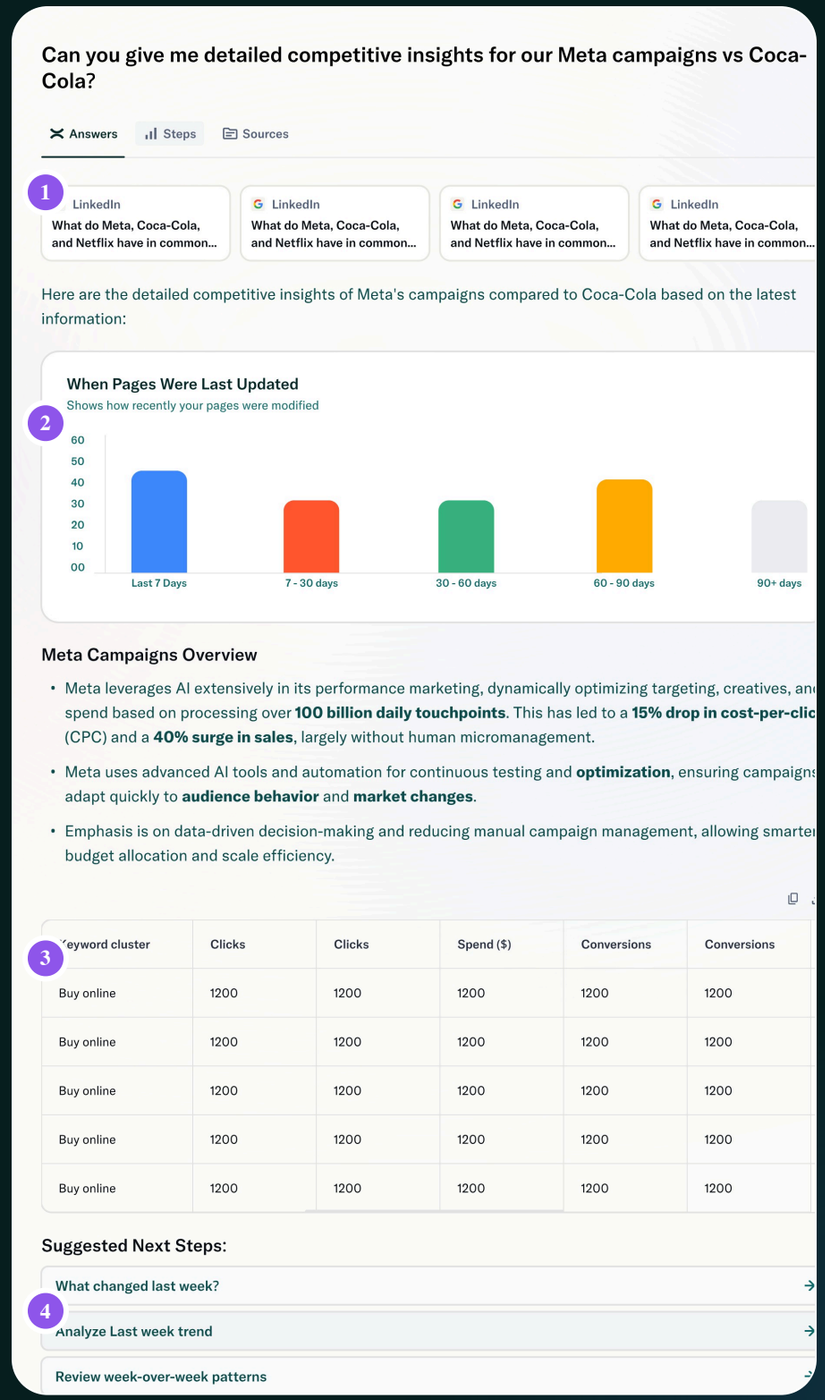
CONTEXT GOES IN WITH THE PROMPT: ACCOUNT, SOURCE OF TRUTH, FILES

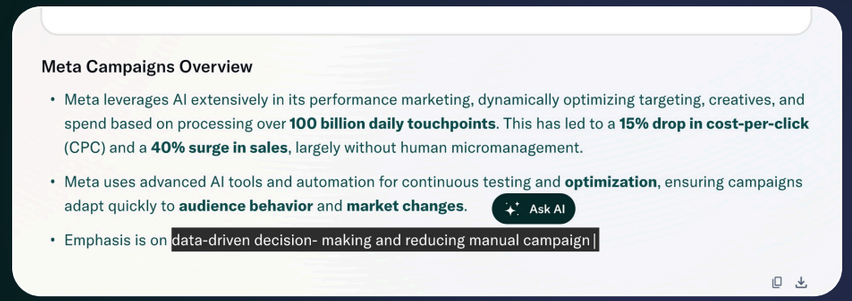
Intent · type, voice, attach → **Scope** · accounts, sources → **Plan** · thoughts, tools → **Answer** · sources, chart, table → **Continue** · next steps, Ask AI → **Recover** · raise an issue

07 FINAL DESIGN

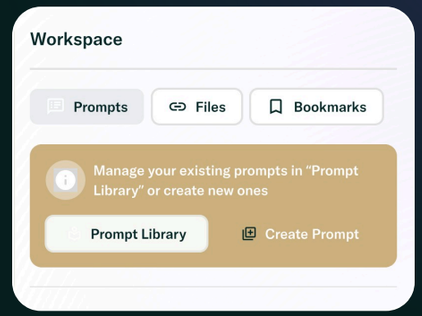
An answer built to be checked, then reused.

1. **Sources first**, above the conclusion
2. **Show, then say**: a chart plus a plain-language overview
3. **Detail on demand**, with copy and download
4. **Next steps** that carry context forward

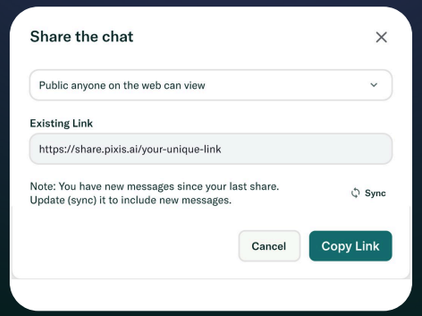




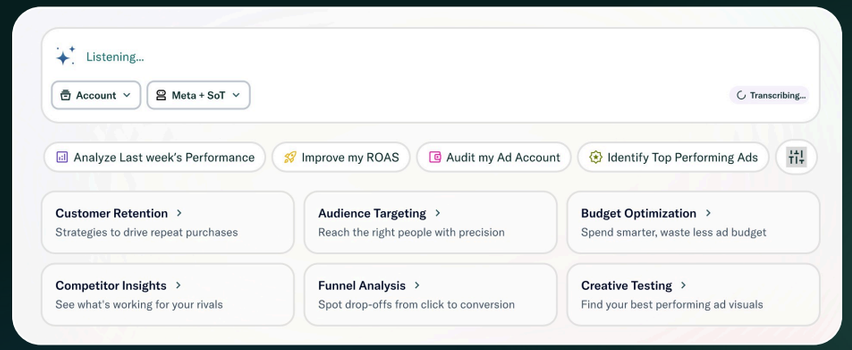
HIGHLIGHT ANY LINE TO ASK AI



PROMPT LIBRARY



SHARE WITH SYNC



VOICE INPUT, WITH LIVE TRANSCRIPTION

A system, a handoff, and a feedback loop for AI answers.

HOW IT MOVED TO BUILD

Low-fi → Mid-fi → Tech validation → Hi-fi → Handover → Dev QA → Iterate

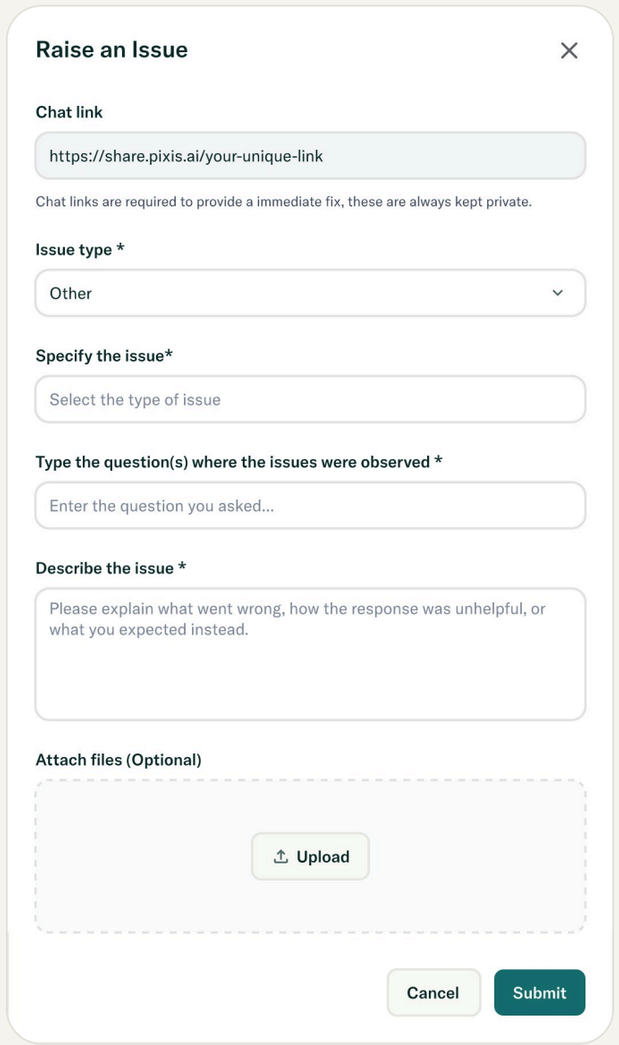
BUILT ON THE PIXIS DESIGN SYSTEM

F60 F40 F10 S20 B10 ACCENT

33 screens and states

1 shared token set

8 modal patterns



RAISE AN ISSUE: A BUG REPORT FOR AI ANSWERS

Chat link required, kept private
Engineering gets the exact conversation to reproduce

The question that failed
Pinpoints the turn where the answer went wrong

What you expected instead
Captures what the user wanted, not just that it failed

Every wrong answer becomes a test case for the AI team.

09 IMPACT AND WHAT'S NEXT

From a half-hour dig to a conversation you can audit.

30×
faster insight-to-decision

70%
less time on campaign audits

100+
team hours saved every month

AS REPORTED BY PIXIS ON PIXIS.AI/PRODUCTS/PRISM, 2026

WHAT I TRACK

Time to a stated reason	**vs 30 min**
Steps / Sources opened	**per answer**
Follow-up depth	**per session**
Issues raised	**per 100 chats**
Prompt reuse via auto-run	**weekly**

WHAT I LEARNED

For expert users, being checkable is part of the answer. They already had answers from tools, ChatGPT and colleagues. What they needed was an answer they could defend in front of a client.

WHAT'S NEXT

Claim-level citations: link every number to its source row

Freshness and conflict states for when the sources disagree

From answer to action: approvals, impact preview, undo

Dashboards show what happened. Stellar explains why.

THANKS FOR READING